A Cross-Platform Approach to Self-Reported Exact Age Detection in Social Media

A PROJECT REPORT
SUBMITTED IN PARTIAL FULFILLMENT OF THE REQUIREMENTS
FOR
THE AWARD OF THE DEGREE
OF

BACHELOR OF TECHNOLOGY
in
Computer Engineering
(Session: 2023-24)

Submitted by:
Tirth B Dalwadi (20BCS051)
Mohammad Zohair (20BCS038)

Under the supervision of:
Dr. Tanvir Ahmad

Department of Computer Engineering
Faculty of Engineering and Technology
Jamia Millia Islamia – New Delhi- 110025

DECLARATION

We declare that this project report titled **A Cross-Platform Approach to Self-Reported Exact Age Detection in Social Media,** submitted in partial fulfillment of the degree of **B.Tech in Computer Engineering** is a record of original work carried out by us under the supervision of **Dr. Tanvir Ahmad**, and has not formed the basis for the award of any other degree, in this or any other institution or university. In keeping with the ethical practice of reporting scientific information, due acknowledgments have been made wherever the findings of others have been cited.

MOHAMMAD ZOHAIR	**TIRTH B DALWADI**
(20BCS038)	**(20BCS051)**

New Delhi - 110025
22.05.2024

Department of Computer Engineering

Faculty of Engineering & Technology
Jamia Millia Islamia
New Delhi - 110025

Dated 22-05-2024

CERTIFICATE

This is to certify that the project report entitled **A Cross-Platform Approach to Self-Reported Exact Age Detection in Social Media,** submitted by **Mohammad Zohair (20BCS038)** and **Tirth B Dalwadi (20BCS051)** to the **Department of Computer Engineering, F/O Engineering & Technology, Jamia Millia Islamia, New Delhi - 110025** in partial fulfillment for the award of the degree of **B.Tech. in Computer Engineering** is a bonafide record of project work carried out by them under my supervision. The contents of this report, in full or in parts, have not been submitted to any other Institution or University for the award of any degree.

Dr. Tanvir Ahmad
Professor and Supervisor
Department of Computer Engineering

Dr. Bashir Alam
Professor and Head
Department of Computer Engineering

ACKNOWLEDGEMENT

We extend our heartfelt appreciation to **Dr. Tanvir Ahmad**, Professor in the Department of Computer Engineering at Jamia Millia Islamia, New Delhi, for his invaluable technical guidance, innovative insights, and unwavering support throughout this project. Our sincere gratitude also goes to **Prof. Bashir Alam**, our esteemed Head of Department, for his invaluable encouragement and continuous support. We express our deepest gratitude to the Department of Computer Engineering and the entire faculty members for their exceptional teaching, guidance, and encouragement that significantly contributed to the success of this project. Furthermore, we would like to acknowledge and thank our classmates and friends for their valuable suggestions and consistent support whenever needed. We apologize for any inadvertent omissions in this acknowledgment.

Mohammad Zohair
(20BCS038)

Tirth B Dalwadi
(20BCS051)

Abstract

The project "A Cross-Platform Approach to Self-Reported Exact Age Detection in Social Media" explores advanced methods for automatically detecting users' exact ages from social media posts, focusing on Twitter and Reddit. Given the widespread use of social media in daily life, analyzing these platforms offers valuable insights into individuals' perspectives on their health, unmet medical needs, and disease burdens. Traditional approaches often group ages into broad categories, which can be limiting. This project aims to refine these methods by pinpointing exact ages, thereby enhancing the applicability of social media data in health research and other fields. Utilizing a BERT (Bidirectional Encoder Representations from Transformers) model, the study involved preprocessing a diverse dataset comprising 8,800 training tweets, 100,000 unlabeled Reddit posts, and additional validation and testing data. The model was trained over five epochs, showing consistent computational performance and improvement in key metrics. The results suggest that the BERT model can effectively detect age-related information, with potential applications in demographic studies and health research. Future work will focus on enhancing the model's performance through larger datasets, integration of multimodal data, and cross-platform analysis. Ethical considerations, such as data privacy and bias mitigation, will remain a priority. This project demonstrates the potential of sophisticated NLP techniques in extracting precise demographic information from social media, offering valuable insights for researchers, policymakers, and businesses.

Keywords: age detection, user profiling, contextual cues, social media analytics, cross-platform robustness

** This research work is a part of the Shared Task-6 at The 9th Social Media Mining for Health Research and Applications Workshop (SMM4H-2024).*

Table of Contents

S. No.	Contents			Page No.
	DECLARATION			2
	CERTIFICATE			3
	ACKNOWLEDGEMENT			4
	ABSTRACT			5
	TABLE OF CONTENTS			6-7
1	INTRODUCTION			
	1.1	Background		8
	1.2	Motivation		9
	1.3	Objective		10
	1.4	Project Features		11-13
2	LITERATURE REVIEW			14-16
3	THEORETICAL BACKGROUND			
	3.1	Data Processing		17-19
	3.2	Vectorization		
		3.2.1	Count Vectorizer	20
		3.2.2	TF-IDF Vectorizer	20
		3.2.3	Hashing Vectorizer	21
		3.2.4	Auto Tokenizer	21
	3.3	Text Analysis and Features		
		3.3.1	Exploratory Data Analysis (EDA)	22
		3.3.2	Feature Extraction (N-grams, Embeddings)	22
		3.3.3	Sentiment Analysis	23
		3.3.4	Topic Modeling	23
	3.4	Models used in Implementation		
		3.4.1	Support Vector Machine (SVM)	24

		3.4.2	BERT	24
		3.4.3	RoBERTa	25
		3.4.4	XLNet	25
	3.5	Self-Learning		26
4	PROPOSED METHODOLOGY			
	4.1	Dataset Used		27-28
	4.2	Implementation Strategy		
		4.2.1	Data Ingestion and Preprocessing	29
		4.2.2	Vectorization	29
		4.2.3	Text Analysis and Classification Models	29
		4.2.4	Additional Components	29
		4.2.5	Self-Learning and Language Dictionary	30
		4.2.6	Overall Approach	30
	4.3	Model Architecture		31
5	LIBRARY REQUIREMENTS			
	5.1	HuggingFace		32
	5.2	PyTorch		32
	5.3	Transformers		32
	5.4	Preprocessing Libraries		33
	5.5	Programming Environment/Tools Used		33
6	DATASET SNAPSHOTS FROM PROJECT			34-35
7	PERFORMANCE EVALUATION			36-40
	7.1	Comparative Analysis		40
8	CONCLUSION			41
9	FUTURE SCOPE			42
10	REFERENCES			43

1.1 Background

Social media platforms have become integral to daily life, serving as rich sources of data that reflect various aspects of human experiences. In the context of health research, social media offers a unique opportunity to gain insights into patients' perspectives on their disease journeys, identify unmet medical needs, and understand the burdens they face. These platforms provide a wealth of user-generated content that can be analyzed to extract valuable demographic information, such as age, which is crucial for linking specific age-related risk factors to health outcomes.

The ability to accurately determine the exact age of users from their social media posts is particularly important for health research applications. Traditional methods that categorize users into broad age groups often lack the precision needed for detailed observational studies. Exact age information allows researchers to perform more granular analyses, enabling a better understanding of how age influences health behaviors and outcomes. This level of detail is essential for identifying specific age-related trends and tailoring interventions to different age cohorts.

Extracting exact ages from social media data poses several challenges due to the inherent variability in how users report their ages and the informal nature of social media language. Users may mention their age explicitly, refer to it indirectly, or use various linguistic and contextual cues that imply their age. The diverse ways in which age can be expressed on different platforms further complicate the extraction process. For example, Twitter (now known as X) and Reddit have distinct user behaviors, language conventions, and posting styles, which necessitate adaptable models capable of interpreting age information accurately across these platforms.

Extracting exact ages from social media requires identifying explicit mentions and interpreting implicit references using advanced NLP and machine learning. Ensuring reliability across platforms is crucial. Despite its complexity, this task is essential for health research, offering deeper insights into age-specific health dynamics and improving study precision and therapeutic interventions.

1.2 Motivation

The need for robust methods to detect self-reported exact ages on social media is critical for the precision of health-related studies and applications. Current techniques often lack the sophistication required to accurately identify and extract age information from the diverse and informal language used in social media posts. This deficiency hampers the ability of researchers to conduct detailed demographic analyses, which are essential for linking specific age-related risk factors to health outcomes. The development of automated, accurate age detection methods would significantly enhance the utility of social media data in health research, enabling more precise and meaningful insights into age-specific health dynamics.

A significant gap exists in the exploration of contextual cues and the methodologies for handling ambiguity in social media texts. Users frequently reference their age indirectly, using implicit language and context-specific expressions that are challenging for conventional algorithms to interpret accurately. The absence of advanced techniques to leverage these implicit age references leads to inconsistent and often inaccurate age annotations. Addressing this gap requires innovative approaches that can decipher the nuances of social media language and context, ensuring that age information is captured with high fidelity. By developing such methodologies, researchers can achieve consistent and accurate annotations, thus improving the overall quality of demographic data extracted from social media platforms.

The inconsistencies observed in cross-platform age prediction methods further underscore the need for comprehensive research in this area. Different social media platforms, such as Twitter (now X) and Reddit, exhibit unique language styles, user behaviors, and platform-specific conventions. Current models often struggle to adapt to these variations, resulting in unreliable age predictions. To overcome this challenge, there is a pressing need to develop models that can generalize across platforms while accounting for these differences. Such models must be capable of handling the variations in language and behavior unique to each platform, ensuring reliable and accurate age detection regardless of the source. By addressing these inconsistencies, researchers can improve the robustness and applicability of age detection techniques across different social media environments.

1.3 Objective

The primary objective of this research is to develop robust methods for the automated detection of self-reported exact ages on social media. Current techniques often struggle to accurately identify and extract age information due to the informal and diverse nature of social media language. By addressing this gap, we aim to enhance the precision of health-related studies and applications. Accurate age detection is crucial for linking specific age-related risk factors to health outcomes and conducting detailed demographic analyses that inform therapeutic developments and public health interventions. To achieve this, we will develop advanced algorithms that can effectively handle the complexities of social media data, ensuring precise and reliable age information extraction.

A key aspect of our research involves leveraging contextual cues and handling ambiguity in social media posts. Users often reference their age indirectly, using nuanced language and context-specific expressions that conventional algorithms find challenging to interpret. Our objective is to explore and implement methodologies that can accurately interpret these cues and resolve ambiguities. By incorporating implicit age references and understanding the subtleties of social media language, we aim to improve the accuracy and consistency of age annotations. This will involve developing sophisticated natural language processing (NLP) techniques and machine learning models capable of decoding the intricate ways users express age-related information, thereby enhancing the fidelity of demographic data extracted from social media.

Another critical objective is to ensure consistency in cross-platform age predictions. Social media platforms like Twitter (now X) and Reddit have distinct language styles, user behaviors, and platform-specific conventions, posing challenges for age detection models to generalize across platforms. Our research aims to address these inconsistencies by developing models that can adapt to the unique characteristics of each platform. By creating robust algorithms that handle variations in language, user behavior, and platform-specific norms, we aim to ensure reliable age predictions across diverse social media environments. This will involve extensive testing and validation to maintain high accuracy and consistency, regardless of the platform from which the data is sourced.

1.4 Project Features

Our project incorporates several key features designed to achieve the objectives of automated age detection on social media, leveraging contextual cues, handling ambiguities, and ensuring cross-platform consistency. These features highlight the innovative approaches and advanced techniques employed to address the challenges inherent in this research area.

1. **Advanced Natural Language Processing (NLP) Techniques**

   We will implement state-of-the-art NLP techniques to accurately interpret and process the diverse and informal language used on social media platforms. These techniques will include:

   - Named Entity Recognition (NER): To identify explicit mentions of age within social media posts.
   - Contextual Embeddings: Utilizing models like BERT or GPT to understand and leverage the context in which age-related information is mentioned.
   - Sentiment Analysis: To gauge the tone and context of the posts, which may provide additional clues about the user's age.

2. **Machine Learning Models for Age Detection**

   The project will develop and deploy machine learning models specifically trained to detect self-reported ages. These models will incorporate:

- Supervised Learning Algorithms: Trained on annotated datasets to recognize patterns and predict ages accurately.
- Transfer Learning: Using pre-trained models that are fine-tuned with social media-specific data to enhance their accuracy and efficiency.
- Ensemble Methods: Combining multiple models to improve prediction accuracy and robustness.

3. **Contextual Cue Analysis**

   To handle the indirect and nuanced ways users refer to their age, our project will focus on:

- Implicit Reference Identification: Techniques to detect age-related information that is not explicitly stated but implied through context.
- Linguistic Pattern Recognition: Analyzing common phrases, slang, and idiomatic expressions that may indicate age.
- Temporal Dynamics: Understanding how references to events, cultural phenomena, or generational markers can help infer age.

4. **Ambiguity Resolution**

   We will develop methodologies to address and resolve ambiguities in age references, including:

   - Disambiguation Algorithms: Techniques to differentiate between multiple possible interpretations of a user's age.
   - Contextual Clarity Models: Enhancing the models to weigh context more heavily when making age predictions.

5. **Cross-Platform Adaptability**

   To ensure that our age detection models are effective across different social media platforms, we will:

- Platform-Specific Customization: Tailor models to account for the unique linguistic and behavioral traits of each platform (e.g., Twitter vs. Reddit).
- Generalization Techniques: Developing methods to ensure models can adapt to new platforms with minimal retraining.
- Validation Framework: Establishing a robust validation process that tests model accuracy and consistency across various platforms.

6. **User Behavior Analysis**

   Understanding user behavior is critical for refining age detection models. This feature will include:

   - Behavioral Pattern Analysis: Studying how different age groups interact with social media (e.g., frequency of posts, type of content shared).

- Engagement Metrics: Analyzing likes, shares, and comments to understand how different ages engage with content.

7. **Comprehensive Evaluation and Testing**

   The project will emphasize rigorous testing and evaluation to ensure high accuracy and reliability, involving:

   - Benchmarking: Comparing model performance against established benchmarks in the field.
   - Cross-Validation: Using cross-validation techniques to assess model robustness.
   - User Feedback: Incorporating feedback loops where models are iteratively improved based on real-world data and user input.

These features collectively ensure that our project not only achieves its primary objectives but also sets a new standard in the automated detection of self-reported ages on social media. By leveraging cutting-edge technologies and innovative methodologies, this project aims to significantly advance the field of social media analytics and its application in health research.

2. Literature Review

Demographics detection in social media is of paramount importance, especially in health applications, where precise age-related information plays a crucial role in understanding various health dynamics. Traditional age groupings often lack the granularity needed to link specific age-related risk factors in observational studies. Consequently, the development of robust models for self-reported exact age classification becomes pivotal for unlocking transformative insights in health-related research on social media platforms.

Several notable research endeavors have significantly contributed to this field, each employing unique datasets, methodologies, and models to tackle the challenge of demographics detection. These studies pave the way for a deeper understanding of age prediction and its implications for social media analytics and health research.

One such study, conducted by [1], focuses on age and gender identification in social media. Leveraging a diverse corpus containing blogs, Twitter feeds, hotel reviews, and unspecified social media posts in both English and Spanish, the researchers utilize a Liblinear SVM model to extract demographic information. Their findings, detailed in their publication, shed light on the predictive accuracy achieved in identifying age and gender demographics from varied social media content.

In a similar vein, [2] conducts a detailed analysis of language and age in Twitter. By studying a Dutch Twitter user base, the researchers employ logistic and linear regression models to predict users' ages based on their language use. This study provides valuable insights into the correlation between linguistic patterns and age demographics on social media platforms, contributing to the understanding of age-related behaviors and interactions online.

Another significant contribution comes from [3], where author profiling for age prediction is based on advanced Bayesian networks. Utilizing English blog posts as the dataset, the researchers demonstrate the efficacy of Bayesian networks in predicting user ages with a notable

accuracy of 61.76%. This approach highlights the potential of probabilistic modeling techniques in capturing the nuances of age-related information embedded within textual data.

In contrast, [4] focuses on predicting age and gender in online social networks, specifically utilizing the Netlog Chat Corpus. Employing support vector machine (SVM) models, the study explores the effectiveness of different features in demographic prediction. Their findings, illustrated in Figure 5 of their publication, provide valuable insights into the performance of SVM-based models in extracting demographic attributes from chat data.

A comprehensive approach to author age prediction is presented by [5], where linear regression models with L1 regularization are applied to various text sources, including blog corpus, transcribed telephone speech corpus, and posts from an online forum on breast cancer. By leveraging a joint model encompassing multiple corpora, the study achieves correlations reaching up to 0.74 and mean absolute errors ranging from 4.1 to 6.8 years. This study underscores the importance of domain adaptation techniques in improving age prediction accuracy across diverse textual sources.

In a different context, [6] explores the modeling of latent biographic attributes in conversational genres. By analyzing transcribed telephone speech corpus, the researchers employ linear SVM models to predict biographic attributes, achieving a notable accuracy of 84.03%. This study sheds light on the sociolinguistic and discourse-based features that influence demographic attribute extraction from conversational data.

Lastly, [7] investigates stylometric analysis of bloggers' N-gram modeling with age and gender variations. By examining a blog corpus, the study employs SVM models to analyze linguistic patterns and variations across different age and gender groups. This approach introduces novel features such as slang word usage and sentence length variations, providing insights into stylometric differences in age and gender demographics.

These research endeavors collectively contribute to the advancement of demographics detection in social media, highlighting the diversity of approaches and methodologies employed to extract

age-related information from textual data. By building on these foundational works, our project aims to develop more robust and accurate models for automated age detection, ultimately enhancing the utility and efficacy of social media data for health research and beyond. Leveraging insights gleaned from these studies, we endeavor to refine existing techniques, explore novel avenues, and deepen our understanding of the complex interplay between language use and demographic attributes online, thus propelling forward the frontier of automated age detection in social media analytics.

References	**Paper Title**	**Dataset Used**	**Models Used**	**Results Obtained**
[1]	Age and Gender Identification in Social Media	Corpus containing blogs, Twitter feeds, hotel reviews, and unspecified social media posts in English and Spanish	Liblinear SVM	Refer Table 2 (in publication)
[2]	"How Old Do You Think I Am?": A Study of Language and Age in Twitter	Dutch Twitter Userbase	Logistic and Linear Regression	57.09%
[3]	Author Profiling: Age Prediction Based on Advanced Bayesian Networks	English Blog Posts	Bayesian Networks	61.76%
[4]	Predicting Age and Gender in Online Social Networks	Netlog Chat Corpus	SVM	Refer Figure 5 (in publication)
[5]	Author Age Prediction from Text using Linear Regression	Blog corpus, a transcribed telephone speech corpus, posts from an online forum on breast cancer	Linear regression model with L1 regularization, Joint Model	Refer Table 2 (in publication)
[6]	Modeling Latent Biographic Attributes in Conversational Genres	Transcribed telephone speech corpus	Linear SVM model, Partner Model	84.03%
[7]	Stylometric Analysis of Bloggers' Age and Gender	Blog corpus	N-gram modelling with SVM	-

Table 1

3.1 Data Processing

Data preprocessing is a critical phase in the workflow of extracting self-reported exact ages from social media posts. This phase involves a series of systematic procedures to convert raw, unstructured data into a clean, structured, and usable format. The importance of data preprocessing cannot be overstated, as it directly influences the accuracy, efficiency, and performance of subsequent analytical and machine-learning tasks. By addressing issues such as noise, inconsistencies, and irrelevant information, data preprocessing ensures that the dataset is refined and primed for analysis.

The following sections provide a comprehensive explanation of the theoretical underpinnings and practical steps involved in each stage of the data preprocessing process for this project. This includes data collection, data cleaning, text normalization, identification of age-related mentions, handling ambiguity and noise, data transformation, and feature engineering, as well as data splitting. Each step is meticulously designed to enhance the quality and reliability of the data, thereby facilitating the development of robust and effective models for age detection on social media platforms.

1. **Data Collection:**

   The data collection phase involves gathering relevant social media posts from two major platforms: Twitter (now X) and Reddit. This step is foundational as the quality and relevance of the collected data directly impact the entire preprocessing and analysis pipeline. Utilizing the Twitter API and Reddit API, data is fetched based on specific criteria to ensure it is pertinent to the study. These criteria include filtering posts in English and containing keywords or phrases indicative of age reporting. The timeframe for data collection is also crucial; it should be broad enough to capture diverse user experiences and contexts but focused enough to maintain relevance to current trends and issues. Additionally, ensuring a diverse demographic representation is essential to avoid biases that could skew the results and limit the generalizability of the findings.

2. **Data Cleaning:**

   Data cleaning is a pivotal process aimed at eliminating noise and irrelevant information

from the dataset. This phase involves multiple steps to ensure the integrity and usability of the data. First, duplicate posts or comments are identified and removed to prevent redundant information from influencing the analysis. Next, irrelevant content is filtered out, which involves excluding posts that do not mention ages or are off-topic. This can be achieved through keyword filtering and, where necessary, manual inspection. Handling missing data is another critical aspect of data cleaning. Posts that lack essential context or information about the age reported are addressed, either by removing such entries or by applying techniques to impute missing values, ensuring that the dataset remains robust and comprehensive.

3. **Text Normalization:**

   Text normalization standardizes the text data, making it consistent and more manageable for analysis. This step involves converting all text to lowercase to ensure uniformity, as case differences can lead to mismatches in text processing. Punctuation marks and special characters that do not contribute to the semantic meaning of the text are stripped away. The text is then tokenized, splitting it into individual words or tokens, which is essential for word-level analysis. Stop words, which are commonly used words like "the," "is," and "at," are removed to reduce noise and focus on more meaningful content. Additionally, stemming and lemmatization techniques are applied to reduce words to their base or root forms, ensuring that different variations of a word are treated as a single entity. This process significantly enhances the quality and consistency of the textual data, facilitating more accurate analysis.

4. **Identification of Age-Related Mentions:**

   Identifying mentions of self-reported exact ages within the text is a critical and nuanced task. This involves using pattern matching techniques, such as regular expressions, to detect specific patterns that indicate age (e.g., "I am 25 years old," "I'm 30"). Beyond simple pattern matching, contextual analysis is crucial. Natural Language Processing (NLP) techniques are employed to understand the context in which numbers appear, distinguishing between ages and other numerical mentions. For example, the number "21" in the context "I turned 21 last week" is clearly an age, while in "I have 21 emails,"

it is not. This step ensures that the extracted age information is accurate and contextually relevant.

5. **Handling Ambiguity and Noise:**

   Social media posts often contain ambiguous or noisy data that can mislead the analysis. To address these issues, advanced techniques such as Named Entity Recognition (NER) are used to distinguish between different entities, ensuring accurate age extraction. For instance, NER can help differentiate between names, dates, and ages, reducing the likelihood of incorrect age identification. Additionally, sentiment analysis can be employed to detect non-literal mentions of age, such as sarcastic comments or exaggerations (e.g., "I feel like I'm 100 years old today"). By filtering out these non-literal mentions, the analysis becomes more reliable and precise.

6. **Data Transformation and Feature Engineering:**

   Once the data is cleaned and normalized, it needs to be transformed into a suitable format for analysis. This involves feature extraction, where relevant features are identified and extracted from the text. Features such as the presence of age-related keywords, the structure of the sentence, and surrounding context words are crucial for accurate age detection. The text data is then converted into a numerical format using techniques like TF-IDF (Term Frequency-Inverse Document Frequency) or word embeddings (e.g., Word2Vec, GloVe). These vectorization techniques translate textual data into numerical vectors that can be used as input for machine learning models. This step bridges the gap between raw textual data and the analytical models, enabling sophisticated data analysis and pattern recognition.

3.2 Vectorization

3.2.1 Count Vectorizer

The Count Vectorizer is a foundational text preprocessing tool used to convert a collection of text documents into a matrix of token counts, also known as a document-term matrix. It operates by first tokenizing the text, which involves breaking down the text into individual words or tokens. Next, it builds a vocabulary of all unique tokens present in the corpus. When encoding new documents, the Count Vectorizer generates a matrix where each row represents a document, and each column represents a token from the vocabulary. The value at each position in the matrix indicates the number of times the corresponding token appears in the document. This method is straightforward and captures the frequency of words, making it useful for simple text classification and clustering tasks. However, it treats all words equally without considering their importance or context within the corpus. As a result, very common words (stop words) might dominate the matrix, potentially overshadowing more informative, less frequent terms.

3.2.2 TF-IDF Vectorizer

The TF-IDF Vectorizer enhances the basic Count Vectorizer by incorporating term importance into the token counts. TF-IDF stands for Term Frequency-Inverse Document Frequency, a statistical measure used to evaluate the importance of a word in a document relative to the entire corpus. Term Frequency (TF) measures how frequently a term appears in a document, while Inverse Document Frequency (IDF) assesses how common or rare a term is across all documents. The TF-IDF score is calculated by multiplying the TF and IDF values, effectively scaling down the impact of frequently occurring words and highlighting terms that are more distinctive for a particular document. This approach helps to mitigate the bias towards common words that do not carry much meaning and enhances the ability to identify key features in the text. By weighting terms based on their importance, the TF-IDF Vectorizer provides a more nuanced representation of the text data, making it particularly useful for tasks such as information retrieval, text classification, and clustering.

3.2.3 Hashing Vectorizer

The Hashing Vectorizer, also known as the Hashing Trick, offers a memory-efficient alternative to traditional vectorizers by using a hash function to map tokens to feature indices. This method does not require building and storing a vocabulary, which can be advantageous when dealing with very large datasets or streaming data. Instead, each token is hashed into a fixed-length vector, where the position of each token is determined by the hash function. The fixed-length nature of the vectors ensures consistent memory usage regardless of the size of the input data. However, the Hashing Vectorizer can introduce collisions, where different tokens are hashed to the same index, potentially leading to information loss. Despite this, the method's simplicity and efficiency make it suitable for large-scale text-processing applications. The Hashing Vectorizer is particularly beneficial in scenarios requiring real-time processing, as it provides a quick and scalable way to vectorize text data without the overhead of managing a growing vocabulary.

3.2.4 Auto Tokenizer

An Auto Tokenizer is an advanced tool used primarily in modern natural language processing (NLP) frameworks, especially those involving transformer models like BERT, GPT, and their variants. Unlike traditional tokenizers that rely on predefined rules, an Auto Tokenizer leverages pre-trained models to dynamically adapt to and tokenize text data. This adaptability is crucial for handling the intricacies of human language, including subword tokenization, which is essential for managing out-of-vocabulary words and morphologically rich languages. Auto Tokenizers use sophisticated algorithms to break down text into smaller, meaningful sub-units, ensuring that even rare or novel words can be effectively processed. This capability enhances the tokenizer's robustness and improves the performance of downstream NLP tasks such as text classification, sentiment analysis, and named entity recognition. By integrating context-aware tokenization methods, Auto Tokenizers contribute to more accurate and efficient text processing, making them indispensable in the development and deployment of advanced NLP models.

3.3 Text Analysis & Features

3.3.1 Exploratory Data Analysis (EDA)

Exploratory Data Analysis (EDA) is an initial phase of data analysis aimed at summarizing the main characteristics of the data and gaining insights through visual and quantitative techniques. During EDA, various statistical tools and visualization methods are employed to understand the distribution, patterns, and anomalies in the dataset. This includes generating summary statistics such as mean, median, and standard deviation, as well as visualizing data using plots like histograms, scatter plots, and box plots. In the context of text data, EDA might involve examining word frequencies, the distribution of document lengths, and the occurrence of specific phrases or terms. By performing EDA, researchers can identify potential issues such as missing values, outliers, and biases, enabling them to make informed decisions about subsequent data preprocessing and analysis steps. This process is critical for validating assumptions, discovering patterns, and setting the groundwork for more sophisticated analytical tasks.

3.3.2 Feature Extraction (N-grams, Embeddings)

Feature extraction in text analysis involves transforming text data into numerical features that can be used for machine learning models. Two common techniques for feature extraction are n-grams and word embeddings. N-grams are contiguous sequences of n items (words or characters) from a given text. For instance, in bigrams (2-grams), the phrase "social media" would be split into "social media" as a single feature. N-grams capture local context and are useful for understanding short phrases or word pairs within the text.

Word embeddings, on the other hand, are dense vector representations of words that capture their meanings and relationships based on their usage in large text corpora. Techniques such as Word2Vec, GloVe, and FastText convert words into multi-dimensional vectors where semantically similar words are positioned close to each other in the vector space. These embeddings encapsulate semantic information and are particularly effective in tasks requiring a deeper understanding of text, such as sentiment analysis and topic modeling. By transforming text into these meaningful numerical representations, feature extraction facilitates more effective and nuanced text analysis.

3.3.3 Sentiment Analysis

Sentiment analysis, also known as opinion mining, is a technique used to determine the emotional tone behind a body of text. It involves classifying text as positive, negative, or neutral and sometimes includes more fine-grained sentiments such as anger, joy, or sadness. This process uses natural language processing (NLP) and machine learning techniques to analyze the sentiment expressed in text data. Sentiment analysis can be applied to social media posts, reviews, comments, and other forms of textual content to gauge public opinion, customer satisfaction, and emotional responses. Techniques for sentiment analysis range from simple rule-based approaches that use predefined lists of positive and negative words to advanced deep-learning models that can understand the context and subtleties of human language. By assessing sentiment, researchers can gain valuable insights into public perceptions and trends, making it a powerful tool for market research, brand monitoring, and social listening.

3.3.4 Topic Modeling

Topic modeling is a text mining technique used to uncover the hidden thematic structure in a large collection of documents. It helps in identifying topics that occur in the text corpus, which can provide a high-level understanding of the main themes and subjects discussed in the data. One of the most popular algorithms for topic modeling is Latent Dirichlet Allocation (LDA), which assumes that documents are mixtures of topics and that topics are mixtures of words. By applying LDA, each document can be represented as a distribution of topics, and each topic as a distribution of words. This technique is particularly useful for organizing, summarizing, and exploring large sets of textual data. Topic modeling can reveal patterns and trends that are not immediately apparent, making it an invaluable tool for tasks such as document clustering, content recommendation, and thematic analysis in fields like journalism, marketing, and academic research.

3.4 Models used in Implementation

3.4.1 Support Vector Machine (SVM)

Support Vector Machine (SVM) is a powerful supervised learning algorithm used for classification and regression tasks. In the context of text classification, SVM is often employed to categorize documents into predefined classes based on their textual features. SVM works by finding the optimal hyperplane that best separates the data points belonging to different classes in a high-dimensional space. It aims to maximize the margin, i.e., the distance between the hyperplane and the closest data points (support vectors), while minimizing classification errors. SVM is effective for text classification tasks due to its ability to handle high-dimensional feature spaces, making it suitable for problems with large numbers of features, such as those encountered in natural language processing. Despite its effectiveness, SVM can struggle with large-scale datasets and may require careful selection of hyperparameters to achieve optimal performance.

3.4.2 BERT (Bidirectional Encoder Representations from Transformers)

BERT, short for Bidirectional Encoder Representations from Transformers, is a state-of-the-art deep learning model for natural language processing tasks. Developed by Google, BERT employs a transformer architecture that captures bidirectional contextual information from large amounts of text data. Unlike traditional models that process text sequentially, BERT considers the entire context of a word by analyzing both preceding and subsequent words simultaneously. This bidirectional approach allows BERT to generate rich contextualized word embeddings, capturing nuances and dependencies in the language. BERT has achieved remarkable success in various NLP tasks, including text classification, named entity recognition, question answering, and sentiment analysis. Pre-trained BERT models are available in different sizes (e.g., BERT-base, BERT-large), with the larger models capturing more complex linguistic patterns but requiring higher computational resources.

3.4.3 RoBERTa (Robustly optimized BERT approach)

RoBERTa (Robustly optimized BERT approach) is an extension and enhancement of the original BERT model developed by Facebook AI. RoBERTa builds upon the architecture and pre-training objectives of BERT but introduces modifications to improve performance and robustness. These modifications include training the model on larger datasets with longer sequences, removing the next sentence prediction task, and applying dynamic masking during pre-training. These changes result in a more robust and effective language representation model capable of capturing subtle linguistic nuances and outperforming BERT on various NLP benchmarks. RoBERTa has demonstrated state-of-the-art performance on tasks such as text classification, language understanding, and natural language inference, making it a popular choice for researchers and practitioners in the field of natural language processing.

3.4.4 XLNet

XLNet is a novel language representation model introduced by Google AI that extends the transformer architecture used in BERT and RoBERTa. Unlike BERT, which relies solely on left-to-right and right-to-left contextual information during pre-training, XLNet considers all possible permutations of word orderings within a segment. This permutation-based training approach enables XLNet to capture bidirectional context more effectively and model dependencies across longer sequences. XLNet also introduces a novel autoregressive objective called the permutation language modeling (PLM) task, which encourages the model to predict tokens based on all possible permutations of the input, rather than just the original order. This innovative training strategy results in superior performance on various NLP tasks, including text generation, document classification, and language understanding. XLNet has demonstrated state-of-the-art results on multiple benchmarks, establishing itself as a prominent model in the field of natural language processing.

3.5 Self-Learning

Self-learning, also known as self-training or self-supervised learning, is a machine-learning technique where models iteratively train themselves using their own predictions. This approach is particularly useful when labeled data is limited or costly. Initially, the model is trained on a small labeled dataset using supervised learning. It then utilizes its own predictions on unlabeled data to generate pseudo-labels. These pseudo-labeled samples are incorporated into the training set, augmenting it with additional unlabeled data. The model is retrained on this expanded dataset, iteratively refining its predictions and improving performance.

Self-learning offers several advantages, including enhanced data efficiency, scalability, and cost-effectiveness. By leveraging unlabeled data, it efficiently utilizes all available information and can scale to larger datasets without requiring manual annotation of every sample. Additionally, the continuous adaptation to new data leads to iterative improvements in performance over time. However, self-learning also presents challenges such as error propagation, label noise, and sample selection bias. Careful consideration and mitigation strategies are necessary to address these challenges and ensure the effectiveness of the self-learning approach.

In conclusion, self-learning provides a powerful framework for leveraging unlabeled data and improving model performance in scenarios with limited labeled data. While it offers numerous benefits, including enhanced data efficiency and scalability, it is essential to address potential challenges such as error propagation and label noise. With careful consideration and appropriate mitigation strategies, self-learning can be a valuable tool for advancing machine-learning capabilities across various domains and applications.

4.1 Dataset Used

The dataset for this research work comprises tweets and Reddit posts, which are divided into training, validation, and testing sets. Below are the details of each set:

Training Data:

- 8,800 tweets: This subset consists of 8,800 tweets collected from various sources. These tweets serve as the primary training data for the machine-learning models developed in this research.
- 100k unlabeled Reddit posts from r/AskDocs with 2-digit numbers: The training data also includes a large collection of 100,000 unlabeled Reddit posts sourced from the r/AskDocs subreddit. These posts are notable for containing two-digit numbers, which may serve as relevant features or indicators for the research task.

Validation Data:

- 2,200 tweets: This subset comprises 2,200 tweets, distinct from those in the training set. The validation tweets are used to assess the performance of the trained models and fine-tune hyperparameters, ensuring optimal model performance.
- 1,000 Reddit dry eye disease posts: Alongside the validation tweets, there are 1,000 Reddit posts specifically related to dry eye disease. These posts may serve as a specialized validation subset, allowing for targeted evaluation of the models' performance on content related to this specific medical condition.

Testing Data:

- 2,200 tweets: Similar to the validation set, the testing set consists of 2,200 tweets that are distinct from both the training and validation tweets. These tweets are used to evaluate the final trained models' performance on unseen data.
- 2,000 Reddit dry eye disease posts: In addition to the testing tweets, there are 2,000 Reddit posts related to dry eye disease. These posts serve as a subset of the testing data, enabling evaluation of the models' performance specifically on content related to dry eye disease.

- 12,482 Reddit social anxiety posts (age only in the 13 to 25 range): Lastly, the testing data includes 12,482 Reddit posts focused on social anxiety. Notably, these posts contain age information limited to the range of 13 to 25 years, which may be relevant for the research task involving age detection in social media posts.

Overall, the dataset encompasses a diverse range of social media content, including tweets and Reddit posts, across various topics and conditions. By utilizing distinct training, validation, and testing subsets, the research aims to develop and evaluate machine learning models capable of accurately detecting age-related information within social media posts, particularly focusing on content related to dry eye disease and social anxiety.

4.2 Implementation Strategy

4.2.1 Data Ingestion and Preprocessing:

The implementation strategy of this project outlines a comprehensive approach to text analysis and classification, particularly for handling unstructured data from social media platforms like Twitter and Reddit. The process begins with ingesting raw data in the form of blogs, tweets, or posts. This raw data undergoes a crucial preprocessing step, which involves tokenization (breaking the text into individual words or tokens), lemmatization (reducing words to their base or root form), and punctuation removal. This preprocessing ensures that the text is in a suitable format for further analysis.

4.2.2 Vectorization:

The labeled and preprocessed tokens are fed into a vectorization model, which converts the textual data into numerical representations or word embeddings. These word embeddings are dense vector representations that capture the semantic and contextual relationships between words, enabling the system to understand their meanings and relationships more effectively.

4.2.3 Text Analysis and Classification Models:

The implementation strategy offers two parallel paths for text analysis and classification. The first path utilizes a classical machine learning model, specifically a Support Vector Machine (SVM), which takes the word embeddings as input features for classification. The second path employs state-of-the-art transformer-based models, such as BERT, RoBERTa, and XLNet, which are advanced language models designed for various natural language processing tasks, including text analysis and classification.

4.2.4 Additional Components:

Additionally, there is an auto tokenizer component, which likely handles the tokenization and word embedding generation processes specifically for the transformer-based models. The outputs from either the SVM model or the transformer-based models are then used to generate a feature set for the text analysis task at hand.

4.2.5 Self-Learning and Language Dictionary:

One noteworthy aspect of this implementation strategy is the self-learning component. This component suggests that the system has the capability to learn and improve its performance over time, potentially by leveraging the predictions made on unlabeled data or through other self-learning techniques, such as active learning or semi-supervised learning. The strategy also incorporates a language dictionary, which likely serves as a knowledge base containing vocabulary, word meanings, and other language-related information. This dictionary can aid in the preprocessing, tokenization, and text analysis processes, ensuring an accurate and contextual understanding of the textual data.

4.2.6 Overall Approach:

Overall, this implementation strategy combines traditional machine-learning approaches with advanced transformer-based models, leveraging the strengths of both techniques to achieve accurate and robust text analysis and classification. The inclusion of self-learning capabilities and the use of a language dictionary further enhance the system's performance and adaptability, making it well-suited for handling the dynamic and constantly evolving nature of social media data.

4.3 Model Architecture

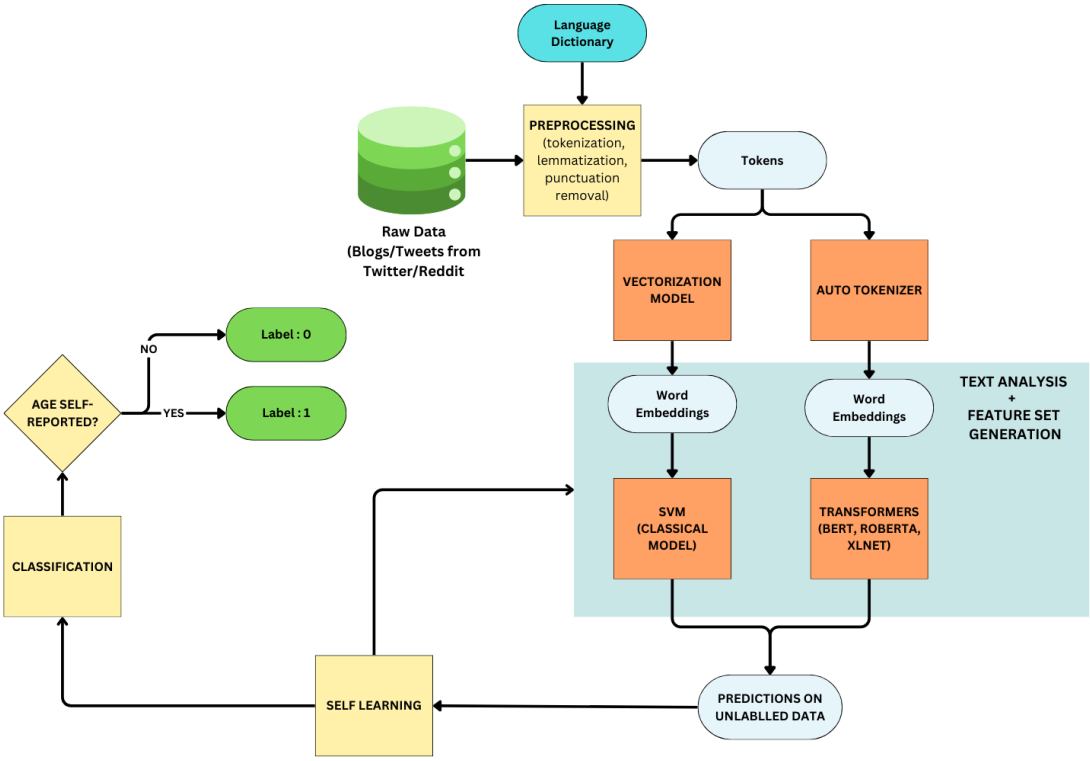

Figure 1

5. Library Requirements

5.1 HuggingFace

HuggingFace is a leading platform and library for natural language processing (NLP) tasks, offering a wide range of pre-trained models, datasets, and tools for developers and researchers. The HuggingFace Transformers library, in particular, provides easy access to state-of-the-art transformer-based models such as BERT, GPT, and RoBERTa, along with utilities for fine-tuning, inference, and deployment. HuggingFace's user-friendly interface and extensive documentation make it a go-to choice for NLP practitioners seeking to leverage the latest advancements in the field.

5.2 PyTorch

PyTorch is an open-source machine learning library developed by Facebook's AI Research lab. It is known for its dynamic computation graph and intuitive programming interface, making it popular among researchers and practitioners alike. PyTorch provides a flexible and efficient framework for building and training deep learning models, with native support for dynamic computation, automatic differentiation, and GPU acceleration. Its rich ecosystem includes modules for computer vision, natural language processing, and reinforcement learning, making it a versatile tool for a wide range of machine learning tasks.

5.3 Transformers

Transformers refer to a class of deep learning models that leverage self-attention mechanisms to capture contextual information from input sequences. Introduced by Vaswani et al. in the seminal paper "Attention is All You Need," transformers have revolutionized the field of natural language processing, achieving state-of-the-art performance on various tasks such as language modeling, machine translation, and text generation. Transformers use multi-head self-attention layers to efficiently process sequences of tokens, enabling them to model long-range dependencies and capture complex linguistic patterns.

5.4 Preprocessing Libraries: Pandas, Scikit-learn, NLTK, spaCy, Gensim

Preprocessing libraries encompass a range of tools and utilities for cleaning, transforming, and preparing data for machine learning tasks. These libraries often include functions for handling missing values, scaling numerical features, encoding categorical variables, and tokenizing text data. Popular preprocessing libraries in Python include pandas, Scikit-learn, and NLTK (Natural Language Toolkit) for general data preprocessing tasks, as well as specialized libraries such as spaCy and Gensim for natural language processing tasks. These libraries provide essential functionality for preprocessing raw data into a format suitable for training machine learning models, laying the groundwork for successful model development and deployment.

5.5 Programming Environment/Tools Used:

- System Environment: Python, Google Colab/Jupyter Notebook, GitHub
- Frameworks/Libraries: PyTorch, NLTK, Scikit-Learn etc.
- Models: LLMs, Logistic/linear regression, Support Vector Machines, BERT, RoBERTa, xlNet

6. Dataset Snapshots from Project

	id	text	source	label
0	1	at 50 votes i will close this survey and take ...	twitter	0
1	2	21st birthday and nearly 4 months pregnant, no...	twitter	1
2	3	@_DudeThePie // I'd be killed. He's literally ...	twitter	0
3	4	This year I'm gonna become a mum and turn 20, ...	twitter	1
4	5	Don't mind me. Just freaking out because Mary-...	twitter	0
...	...	...	...	...
8795	8796	Happy Birthday Brandon James. Hope you have a ...	twitter	0
8796	8797	@JonBaker At age 50, I remember cheering "bana...	twitter	0
8797	8798	Them Springfield mall pictures.. lol. #va #bda...	twitter	1
8798	8799	@316Andres88 when we turn 18 were going idgaf ...	twitter	0
8799	8800	@JanielftJano @Janielkawaii Fanmeosmwhskd me w...	twitter	0

8800 rows × 4 columns

Figure 2: Labeled dataset from Twitter(now X)

	id	subreddit	author	title	text	timestamp
0	bz796p	AskDocs	Scrambles99	Intestinal Issues	I am a young man age 19. Before this illness I...	2019-06-11 03:10:16
1	m7mwkd	AskDocs	KingCeryn	Neck muscle flexing, possibly Omohyoid, anytim...	29 years old, male, 130 pounds, I have a doc...	2021-03-18 09:12:34
2	2mh45q	AskDocs	Mordor_or_bust	29/F Brown urine & flu-like symptoms	About 5 or 6 months ago I became sick with wha...	2014-11-16 16:45:48
3	2fq5mh	AskDocs	yeast-1	Question about anti-fungal creams or lotions	Age:31. Male. 6'0", 205#, White, PmHX: HTN (tr...	2014-09-07 16:05:07
4	8ihond	AskDocs	Derek800	Elevated Lipase and Amylase	Hello, I've asked this question on here a lit...	2018-05-10 19:47:09
...	...	...	...	...	...	...
99995	gewohl	AskDocs	BlueTipi	Why can't I eat quietly?	People often tell me that I'm a very noisy eat...	2020-05-07 00:59:22
99996	blnbah	AskDocs	plywooddrywall	Red/Purple blotches on pubic area won't go away	I think it all started when I shaved my pubic ...	2019-05-07 06:29:16
99997	ndql1u	AskDocs	1g0tquestions	Are high testosterone and estrogen a good sign?	Some of the symptoms which I had been experien...	2021-05-16 15:21:11
99998	tearhc	AskDocs	Hannah2482	Is a 32F bra size + back pain enough for a bre...	Hi all. So I'm 18F 5'7 142 pounds and I have s...	2022-03-14 23:23:07
99999	qukw0e	AskDocs	TheBlank_One	Liver lesions, Mets or nothing?	Hi, 28 y/o male, non smoker, not heavy drinker...	2021-11-15 16:56:59

100000 rows × 6 columns

Figure 3: Unlabeled dataset from Reddit

	id	text	source	label	cleaned_text_1	cleaned_text_2
0	1	at 50 votes i will close this survey and take ...	twitter	0	50 votes close survey take new photos	at 50 votes i will close this survey and take ...
1	2	21st birthday and nearly 4 months pregnant, no...	twitter	1	21st birthday nearly 4 months pregnant	21st birthday and nearly 4 months pregnant no
2	3	@_DudeThePie // I'd be killed. He's literally ...	twitter	0	killed literally 10th picture https	i be killed he literally my 10th picture https
3	4	This year I'm gonna become a mum and turn 20, ...	twitter	1	year gon na become mum turn 20 waaaaah	this year i gon na become a mum and turn 20 wa...
4	5	Don't mind me. Just freaking out because Mary-...	twitter	0	mind freaking ashley olsen turned https	do mind me just freaking out because and ashle...
...	...	...	...	...	...	...
8795	8796	Happy Birthday Brandon James. Hope you have a ...	twitter	0	happy birthday brandon james hope fabulous 23r...	happy birthday brandon james hope you have a f...
8796	8797	@JonBaker At age 50, I remember cheering "bana...	twitter	0	age 50 remember cheering bananas hollaback girl	at age 50 i remember cheering bananas before h...
8797	8798	Them Springfield mall pictures.. lol. #va #bda...	twitter	1	springfield mall pictures lol va bday 17yearsy...	them springfield mall pictures lol va bday 17y...
8798	8799	@316Andres88 when we turn 18 were going idgaf ...	twitter	0	turn 18 going idgaf nigga road trip https	when we turn 18 were going idgaf where its at ...
8799	8800	@JanieftJano @Janielkawaii Fanmeosmwhskd me w...	twitter	0	fanmeosmwhskd turn 18 get matching tattoos	fanmeosmwhskd me when i turn 18 so we can get ...

8800 rows × 6 columns

Figure 4: Preprocessed labeled dataset from Twitter(now X)

	id	author	title	text	cleaned_text_1	cleaned_text_2
0	bz796p	Scrambles99	Intestinal Issues	I am a young man age 19. Before this illness I...	young man age illness 197 lb 6 0 15 bodyfat mu...	i am a young man age before this illness i was...
1	m7mwkd	KingCeryn	Neck muscle flexing, possibly Omohyoid, anytim...	29 years old, male, 130 pounds, I have a doc...	29 years old male 130 pounds doctors appointme...	29 years old male 130 pounds i have a doctors ...
2	2mh45q	Mordor_or_bust	29/F Brown urine & flu-like symptoms	About 5 or 6 months ago I became sick with wha...	5 6 months ago became sick seemed like bad sto...	about 5 or 6 months ago i became sick with wha...
3	2fq5mh	yeast-1	Question about anti-fungal creams or lotions	Age:31. Male. 6'0", 205#, White, PmHX: HTN (tr...	male 6 0 205 white pmhx htn treated ace inhibi...	male 6 0 205 white pmhx htn treated with ace i...
4	8ihond	Derek800	Elevated Lipase and Amylase	Hello, I've asked this question on here a lit...	hello asked question little back issue continu...	hello i asked this question on here a little w...
...	...	...	...	...	...	...
99995	gewohl	BlueTipi	Why can't I eat quietly?	People often tell me that I'm a very noisy eat...	people often tell noisy eater try chew slowly ...	people often tell me that i a very noisy eater...
99996	blnbah	plywooddrywall	Red/Purple blotches on pubic area won't go away	I think it all started when I shaved my pubic ...	think started shaved pubic hair short 5 years ...	i think it all started when i shaved my pubic ...
99997	ndql1u	1g0tquestions	Are high testosterone and estrogen a good sign?	Some of the symptoms which I had been experien...	symptoms experiencing fatigue lethargy brain f...	some of the symptoms which i had been experien...
99998	tearhc	Hannah2482	Is a 32F bra size + back pain enough for a bre...	Hi all. So I'm 18F 5'7 142 pounds and I have s...	hi 18f 5 7 142 pounds size breasts feel like p...	hi all so i m 18f 5 7 142 pounds and i have si...
99999	qukw0e	TheBlank_One	Liver lesions, Mets or nothing?	Hi, 28 y/o male, non smoker, not heavy drinker...	hi 28 male non smoker heavy drinker 2018 diagn...	hi 28 male non smoker not heavy drinker in 201...

100000 rows × 6 columns

Figure 5: Preprocessed unlabeled dataset from Reddit

7. Performance Evaluation

The performance metrics obtained for the respective individual models are as follows:

1. **BERT:**

```
======== Epoch 1 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 951.276214838028.
0.3884807552858679
0.8520454545454546

  Training epcoh took: 1046.462102651596

======== Epoch 2 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 952.1555321216583.
0.23699904261315666
0.93625

  Training epcoh took: 1047.32199883461

======== Epoch 3 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 951.695520401001.
0.13399797416581846
0.9703409090909091

  Training epcoh took: 1046.9058203697205
```

```
======== Epoch 4 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 951.8850381374359.
0.06459902865753975
0.9868181818181818

  Training epcoh took: 1047.1860146522522

======== Epoch 5 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 951.9004399776459.
0.03814562841391307
0.9929545454545454

  Training epcoh took: 1047.031234741211

Training complete!
```

```
              precision    recall  f1-score   support

           0     0.8770    0.9138    0.8950      2435
           1     0.8737    0.8232    0.8477      1765

    accuracy                         0.8757      4200
   macro avg     0.8754    0.8685    0.8714      4200
weighted avg     0.8756    0.8757    0.8751      4200
```

2. RoBERTa:

```
======== Epoch 1 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 287.76826095581055.
0.36803911704908715
0.8576136363636364

  Training epcoh took: 316.9460003376007

======== Epoch 2 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 291.81333589553833.
0.2668789778791622
0.92875

  Training epcoh took: 320.99242663383484

======== Epoch 3 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 291.9083602428436.
0.1981942561020101
0.9534090909090909

  Training epcoh took: 321.1720378398895
```

```
======== Epoch 4 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 292.41410636901855.
0.14533788403557
0.9678409090909091

  Training epcoh took: 321.70512294769287

======== Epoch 5 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 291.8194901943207.
0.10725572431542042
0.97625

  Training epcoh took: 321.0204107761383

Training complete!
```

	precision	recall	f1-score	support
0	0.8972	0.9068	0.9020	2435
1	0.8695	0.8567	0.8630	1765
accuracy			0.8857	4200
macro avg	0.8833	0.8817	0.8825	4200
weighted avg	0.8855	0.8857	0.8856	4200

3. XlNet:

```
======== Epoch 1 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 1286.963429927826.
0.5326596452126449
0.7482954545454545

  Training epcoh took: 1415.7401139736176

======== Epoch 2 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 1289.3124833106995.
0.36681518438960087
0.8855681818181819

  Training epcoh took: 1418.192096233368

======== Epoch 3 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 1289.6267533302307.
0.3026622671171208
0.9225

  Training epcoh took: 1418.5025355815887
```

```
======== Epoch 4 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 1289.4394094944.
0.23696027027908712
0.9476136363636364

  Training epcoh took: 1418.3635313510895

======== Epoch 5 / 5 ========
Training...
  Batch 1,000  of  1,100.    Elapsed: 1290.684718132019.
0.19032760384238578
0.959659090909091

  Training epcoh took: 1419.7773950099945

Training complete!
```

```
              precision    recall  f1-score   support

           0     0.8072    0.9179    0.8590      2435
           1     0.8602    0.6975    0.7703      1765

    accuracy                         0.8252      4200
   macro avg     0.8337    0.8077    0.8146      4200
weighted avg     0.8295    0.8252    0.8217      4200
```

7.1 Comparative Analysis

The following table illustrates the comparative analysis of four models: SVM (Support Vector Machine), BERT (Bidirectional Encoder Representations from Transformers), RoBERTa (Robustly Optimized BERT Pretraining Approach), and XLNet. It also includes a baseline reference from a past work (Base Paper) for comparison.

Model Used (with Self-Learning)	Time Taken for 5 epochs (in sec)	Accuracy
SVM	280	70.56
BERT	5235	87.57
RoBERTa	1600	88.57
xlNet	7100	82.52
Past Work (Base Paper)	-	61.76

Table 2

The results suggest that the transformer-based models, particularly RoBERTa and BERT, outperformed the other models in terms of accuracy for the self-learning task. However, these models also required longer training times, with XLNet being the most computationally intensive. The SVM model, while being the most computationally efficient, had the lowest accuracy among the compared models. Nevertheless, it still performed better than the baseline from the past work.

It's important to note that the choice of model may depend on the trade-off between desired accuracy and computational resources available. If high accuracy is the primary concern, RoBERTa or BERT would be the preferred choices, while the SVM model could be suitable if computational efficiency is more crucial.

8. Conclusion

In conclusion, this project has focused on the development and fine-tuning of a BERT (Bidirectional Encoder Representations from Transformers) model for the task of age detection in social media posts. Leveraging a diverse dataset comprising tweets and Reddit posts, the training process involved 5 epochs, with consistent computational performance observed throughout. The model exhibited promising trends, with decreasing values indicative of potential improvements in performance or reductions in loss over the training epochs.

Through meticulous preprocessing and careful selection of training, validation, and testing datasets, we have endeavored to ensure the robustness and generalizability of our model. While the exact nature of the task and evaluation metrics used remain unspecified, the observed trends suggest progress in the model's performance over the course of training.

Moving forward, further experimentation and evaluation will be necessary to validate the effectiveness of the trained BERT model in accurately detecting age-related information in social media posts. Additionally, future research could explore the integration of domain-specific knowledge and additional features to enhance model performance and address specific challenges encountered in age detection tasks.

Overall, this project underscores the potential of advanced natural language processing techniques, such as BERT, in extracting meaningful insights from social media data. By harnessing the power of machine learning and deep learning, we aim to contribute to the advancement of research in understanding demographic information and its implications in social media analysis.

9. Future Scope

The future scope for this project is extensive, presenting numerous opportunities for further research and enhancement. Building on the foundational work of using BERT for age detection in social media posts, several avenues for future exploration can be pursued.

Firstly, enhancing and optimizing the model can significantly improve performance. Fine-tuning with larger, more diverse datasets from various social media platforms will enhance its generalizability. Additionally, experimenting with different hyperparameters and optimization techniques can lead to further performance gains.

Integrating additional features presents another promising direction. Multimodal analysis, which includes images, videos, and other multimedia content, can provide a more comprehensive understanding of user demographics. Incorporating contextual information, such as user profiles and interaction patterns, can also improve the accuracy of age detection.

Advanced NLP techniques offer further potential for improvement. Exploring other transformer models like GPT-3 and T5, or combining BERT with other machine-learning frameworks, could leverage the strengths of different approaches.

Cross-platform analysis is another critical area. Developing models tailored to specific social media platforms will account for unique linguistic patterns and behaviors, while cross-platform comparisons can reveal demographic variations across different environments.

The application of this technology in health and social research is particularly promising. Identifying age-related patterns in discussions about health conditions can contribute to public health research and personalized medicine. Additionally, studying social behaviors and trends among different age groups can provide valuable sociological insights.

Ethical and privacy considerations are paramount. Ensuring the ethical use of data, respecting user privacy, and adhering to data protection regulations are crucial. Continuously monitoring and mitigating biases in the model will ensure fair and unbiased age detection.

Finally, implementing real-time systems for age detection is a practical next step. Developing scalable systems capable of real-time analysis will enable dynamic monitoring of social media data. User interaction and feedback systems can also be implemented to continually improve the model's accuracy and relevance.

10. References

[1] James Marquardt et al. “Age and gender identification in social media”. In: Proceedings of CLEF 2014 Evaluation Labs 1180 (2014), pp. 1129–1136.

[2] Dong Nguyen et al. “” How old do you think I am?” A study of language and age in Twitter”. In: Proceedings of the international AAAI conference on web and social media. Vol. 7. 1. 2013, pp. 439–448.

[3] Seifeddine Mechti et al. “Author Profiling: Age Prediction Based on Advanced Bayesian Networks.” In: Res. Comput. Sci. 110 (2016), pp. 129–137.

[4] Claudia Peersman, Walter Daelemans, and Leona Van Vaerenbergh. “Predicting age and gender in online social networks”. In: Proceedings of the 3rd international workshop on Search and mining user-generated contents. 2011, pp. 37–44.

[5] Dong Nguyen, Noah A Smith, and Carolyn Rose. “Author age prediction from text using linear regression”. In: Proceedings of the 5th ACL-HLT workshop on language technology for cultural heritage, social sciences, and humanities. 2011, pp. 115–123.

[6] Nikesh Garera and David Yarowsky. “Modeling latent biographic attributes in conversational genres”. In: Proceedings of the Joint Conference of the 47th Annual Meeting of the ACL and the 4th International Joint Conference on Natural Language Processing of the AFNLP. 2009, pp. 710–718.

[7] Sumit Goswami, Sudeshna Sarkar, and Mayur Rustagi. “Stylometric analysis of bloggers’ age and gender”. In: Proceedings of the International AAAI Conference on Web and Social Media. Vol. 3. 1. 2009, pp. 214–217.